

PLENSIA LUKOSI

DATA ANALYST

Dar es Salaam, Tanzania · +255 699 675 407 · plensiapl@gmail.com · linkedin.com/in/plensia-lukosi · github.com/Plensia

Data analyst who turns a vague ask into a checkable question, then answers it in SQL, Power BI and Excel. Four projects published with their method, their bugs and their limits — a 6.36M-row fraud-rule comparison, a containerized ETL and dbt warehouse, a 59-post content audit, and a 4,023-respondent survey dashboard. BSc Computer Science, Google Advanced Data Analytics certified. Open to data analyst roles — remote or Dar es Salaam.

TECHNICAL SKILLS

SQL PostgreSQL, MySQL — joins, aggregations, window functions, views, schema design, data validation and reconciliation

BI & reporting Power BI (DAX, Power Query / M, data modelling, interactive dashboards); Excel (Power Query, PivotTables, VLOOKUP / XLOOKUP, dashboards); Looker Studio

Python & data engineering pandas, Matplotlib, Streamlit; incremental ETL design, dbt modelling, Docker Compose, Git / GitHub

Analysis practice metric definition, data-quality auditing, root-cause analysis, reporting automation, clear written documentation

SELECTED PROJECTS — CODE AND WRITE-UPS ON GITHUB

PaySim Fraud Detection — Python · PostgreSQL · Power BI [Plensia/PaySim-Fraud-Analysis](https://github.com/Plensia/PaySim-Fraud-Analysis)

- Framed “catch fraud without freezing legitimate merchant cash flow” as a rule comparison over 6.36M mobile-money transactions: analysis in pandas, re-expressed as PostgreSQL views, delivered as a Power BI dashboard.
- Showed the incumbent rule was effectively dormant (16 of 8,213 fraud cases, 0.2% recall) and that the obvious fix overcorrects — a flat 200,000 threshold reaches 66.6% recall at 1,192,198 false positives, roughly 218 false alarms per fraud caught.
- Recommended one behavioural signal instead — the sender's balance drained to zero — which catches **97.7% of fraud with zero false positives**; shipped a filterable flagged-accounts view a fraud review team can work from.

Maven Fuzzy Factory — end-to-end data stack — Python · PostgreSQL · dbt · Docker · Streamlit [Plensia/maven-fuzzy-factory-etl](https://github.com/Plensia/maven-fuzzy-factory-etl)

- Built a containerized ETL pipeline loading 1.2M+ e-commerce rows into PostgreSQL, with primary-key validation and incremental loading on a `created_at` watermark, so a rerun extends the warehouse instead of duplicating it.
- Modelled five dbt analytics marts — channel performance, product profitability, refund trends, cohort retention, order value — surfaced through a Streamlit executive dashboard, all orchestrated with Docker Compose.

LinkedIn Content Performance — Excel · *professional work, DataGirls Tanzania* [Plensia/linkedin-content-performance-analysis](https://github.com/Plensia/linkedin-content-performance-analysis)

- Tagged and audited 59 organic posts (Jan–Nov 2024) by topic and length — neither field existed in the export — to separate reach from conversion: event-recap posts drew just 1,209 impressions but a **39.5% engagement rate, over double the next-best topic**.
- Fixed the data-quality defects distorting the charts — inconsistent category casing splitting one topic in two, a “Total” row embedded in the raw table injecting phantom categories — and disproved “shorter is better”: medium-length posts led click-through at 11.4% against 3.6% for the shortest.

Great American Coffee Taste Analysis — Power BI · Power Query (M) [Plensia/great-american-coffee-taste-analysis](https://github.com/Plensia/great-american-coffee-taste-analysis)

- Modelled a 4,023-respondent survey to test stated preference against blind taste results — the same coffee won both (36.7% stated favourite, 3.38 blind score) — and corrected three ordinal charts sorting by response count rather than the survey's canonical sequence, while leaving genuinely nominal categories on popularity sort.

EXPERIENCE

DataGirls Tanzania — Digital Media & Communications Manager Dodoma, Tanzania · Sep 2025 – Present

- Own the LinkedIn content analytics reporting above: maintain the monthly performance workbook, tag posts by topic and format, and report engagement against community growth benchmarks.
- Reconcile platform metadata across Instagram, TikTok and LinkedIn to surface anomalies and hold metric definitions consistent between sources; automated recurring reports and validation checks previously rebuilt by hand each month.

DataGirls Tanzania — Assistant Community Manager Dodoma, Tanzania · Feb 2025 – Aug 2025

- Tracked and reconciled event engagement metrics across platforms, supporting continuous-improvement work that raised satisfaction scores by 15%.

UNDP Tanzania — Student Ambassador (Cohort 3 & extended term) Dar es Salaam · Apr 2025 – Present

- Designed data-informed digital content for university networks and facilitated data-literacy workshops for peer cohorts.

EDUCATION & CERTIFICATIONS

BSc Computer Science — University of Dodoma

Google Advanced Data Analytics — Coursera, Nov 2025 · **Associate Data Analyst** — DataCamp, May 2026 · **Power BI Desktop for Business Intelligence** — Maven Analytics / Udemy, Sep 2025 · **Microsoft PL-300** — Power BI Data Analyst Associate, in progress